



PROJECT NAME :
TITLE: *IBM EMPLOYEE ATTRITION*

PREDICTIVE MODELING

Batch:
DSP 19, Group 3

Submitted by:
Tejinder Singh Wadhwa & Qays Ibrahim Tole


INTRODUCTION:

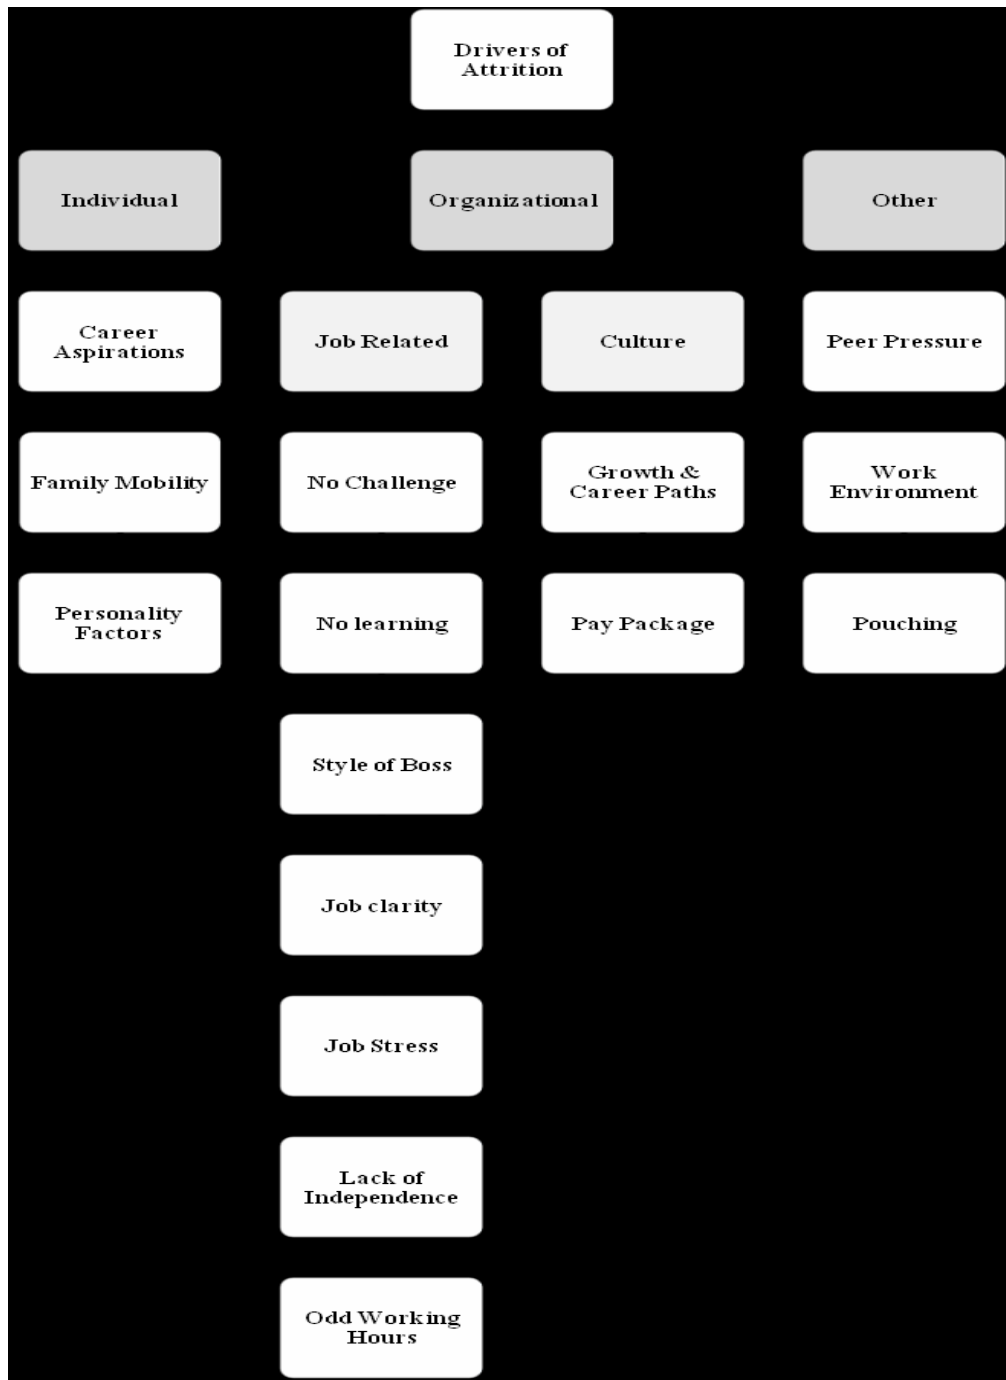
What is Attrition?

Attrition simply means "A reduction in the number of employees through retirement, resignation or death."

Attrition Scenario in India

Almost all the sectors in India are facing attrition, but the reasons and effects are unique to each sector. The attrition rates of different sector are illustrated below:





Cost of Attrition

Employee attrition is a costly dilemma for all organizations. Employee attrition costs 12 to 18 months' salary for each leaving manager or professional.



OBJECTIVE:

OBJECTIVE 1:

Identifying the factors that cause employees to leave the organization and explore important questions such as 'compare average monthly income by education and attrition' or 'is distance from home a case for attrition'?

OBJECTIVE 2:

Build classification model to predict which employee is likely to churn and help the business to devise policies and attract back the right talent.



TOOLS & PACKAGES USED:

Tools:

1. **R STUDIO:** We used this for importing the data, handling outliers, scaling, creating models & their summary as well as creating histogram charts
2. **Tableau:** We have used tableau for creating various pie charts and line graphs for presentation

Packages:

1. **caret:** The **caret** package (short for Classification and Regression Training) contains functions to streamline the model training process for complex regression and classification problems. One of the primary tools in the package is the `train` function which we used it in our coding
2. **e071:** Functions for latent class analysis, short time Fourier transform, fuzzy clustering, support vector machines, shortest path computation, bagged clustering, naive Bayes classifier. We used it for creating SVM models.
3. **ggplot2:** ggplot2 is a system for declaratively creating graphics, based on The Grammar of Graphics. Used for crating box plots, histograms etc.
4. **rminer:** We used this for finding important variables as per SVM model.
5. **randomForest:** Implements Breiman's random forest algorithm (based on Breiman and Cutler's original Fortran code) for classification and regression. We used this package to create Random Forest models



FEATURE SELECTION:

Deleting Variables:

Data file had **35 variables**. We deleted the following variables from the data set for the reasons mentioned below:

Variable Name	Reason for deletion
EmployeeCount	Variable had only 1 level which gives error when included in model creation
EmployeeNumber	Unique number for each employee. No significance in model creation
Over18	Variable had only 1 level which gives error when included in model creation
StandardHours	Variable had only 1 level which gives error when included in model creation

Creating New variable:

If we have a look at the datafile then you come across the below three variables that are related to the employee job satisfaction:

- i. EnvironmentSatisfaction = Employee satisfaction score from 1 to 4 with 1 = Low to 4 = Very High
- ii. JobSatisfaction = Employee Job satisfaction score from 1 to 4 with 1 = Low to 4 = Very High
- iii. RelationshipSatisfaction = Employee Job satisfaction score from 1 to 4 with 1 = Low to 4 = Very High

If you use scores of the above three features in any combination then we can derive how satisfied that employee is with his/her job in the organization

Hence; we thought it would be a nice idea to derive a new variable which will combine the scores of the above three variables and provide us the employee happiness index with his/her job. We named it as "Work_Happiness_Index"

We use the below condition to derive Happiness index of the employees:

- i. If JobSatisfaction ≥ 3 and EnvironmentSatisfaction ≥ 3 AND RelationshipSatisfaction ≥ 3) we categorized such employees "Work_Happiness_Index" as **"Very Happy"**
- ii. If JobSatisfaction ≥ 3 AND (EnvironmentSatisfaction ≥ 3 OR RelationshipSatisfaction ≥ 3) we categorized such employees "Work_Happiness_Index" as **"Happy"**
- iii. If JobSatisfaction ≤ 2 AND (EnvironmentSatisfaction ≥ 3 OR RelationshipSatisfaction ≥ 3) we categorized such employees "Work_Happiness_Index" as **"Sad"**
- iv. If JobSatisfaction ≤ 2 AND EnvironmentSatisfaction ≤ 2 AND RelationshipSatisfaction ≤ 2 we categorized such employees "Work_Happiness_Index" as **"Very Sad"**
- v. Rest Employees who didn't fall in the above 4 categories are classified a "Work_Happiness_Index" as **"Somewhat Happy"**.

After Feature Selection the original data file is left with **32 Variables** only which were used for Model building.



ALGORITHMS USED:

- **Logistic regression**

Logistic regression measures the relationship between the categorical dependent variable and one or more independent variables by estimating probabilities using a logistic function.

- **Decision tree**

Decision Tree algorithm belongs to the family of supervised learning algorithms. Unlike other supervised learning algorithms, decision tree algorithm can be used for solving regression and classification problems too.

The decision tree algorithm tries to solve the problem, by using tree representation. Each internal node of the tree corresponds to an attribute, and each leaf node corresponds to a class label.

Decision Tree Algorithm Pseudocode:

- i. Place the best attribute of the dataset at the root of the tree.
- ii. Split the training set into subsets. Subsets should be made in such a way that each subset contains data with the same value for an attribute.
- iii. Repeat step 1 and step 2 on each subset until you find leaf nodes in all the branches of the tree.

- **Random Forests:**

Random forests or random decision forests are an ensemble learning method for classification, regression and other tasks that operates by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes (classification) or mean prediction (regression) of the individual trees. Random forest is like bootstrapping algorithm with Decision tree model.


- **SVM:**

Support-vector machines (SVMs, also support-vector networks are supervised learning models with associated learning algorithms that analyze data used for classification and regression analysis. Given a set of training examples, each marked as belonging to one or the other of two categories, an SVM training algorithm constructs a hyperplane or set of hyperplanes in a high- or infinite-dimensional space, which can be used for classification, regression, or other tasks like outliers detection


DATA SPLITTING:

- We have made a random split of the ratio **70:30** respectively using the sample() function of R.
- Where in 70% of the data represents training data with **1029 observations and 32 variables**. We named this file as **"trainData.csv"**.
- And 30% of the data represents the testing data with **441 observation and 32 variables**. We named this file as **"testData.csv"**.
- Below graphs shows the percentage split of "Yes" and "No" of our target variable "Attrition" in Train and Test data is almost same as it was in the original data file. This implies that it is a **good data split**.
- As you can see from the below pie charts the ratio of count of "No" Attrition is way more than the "Yes" counts. It's almost 83:17 in the original file itself. This could lead to creation of models little bias towards predicting "No" Attrition more accurately than "Yes".


org data




traindata



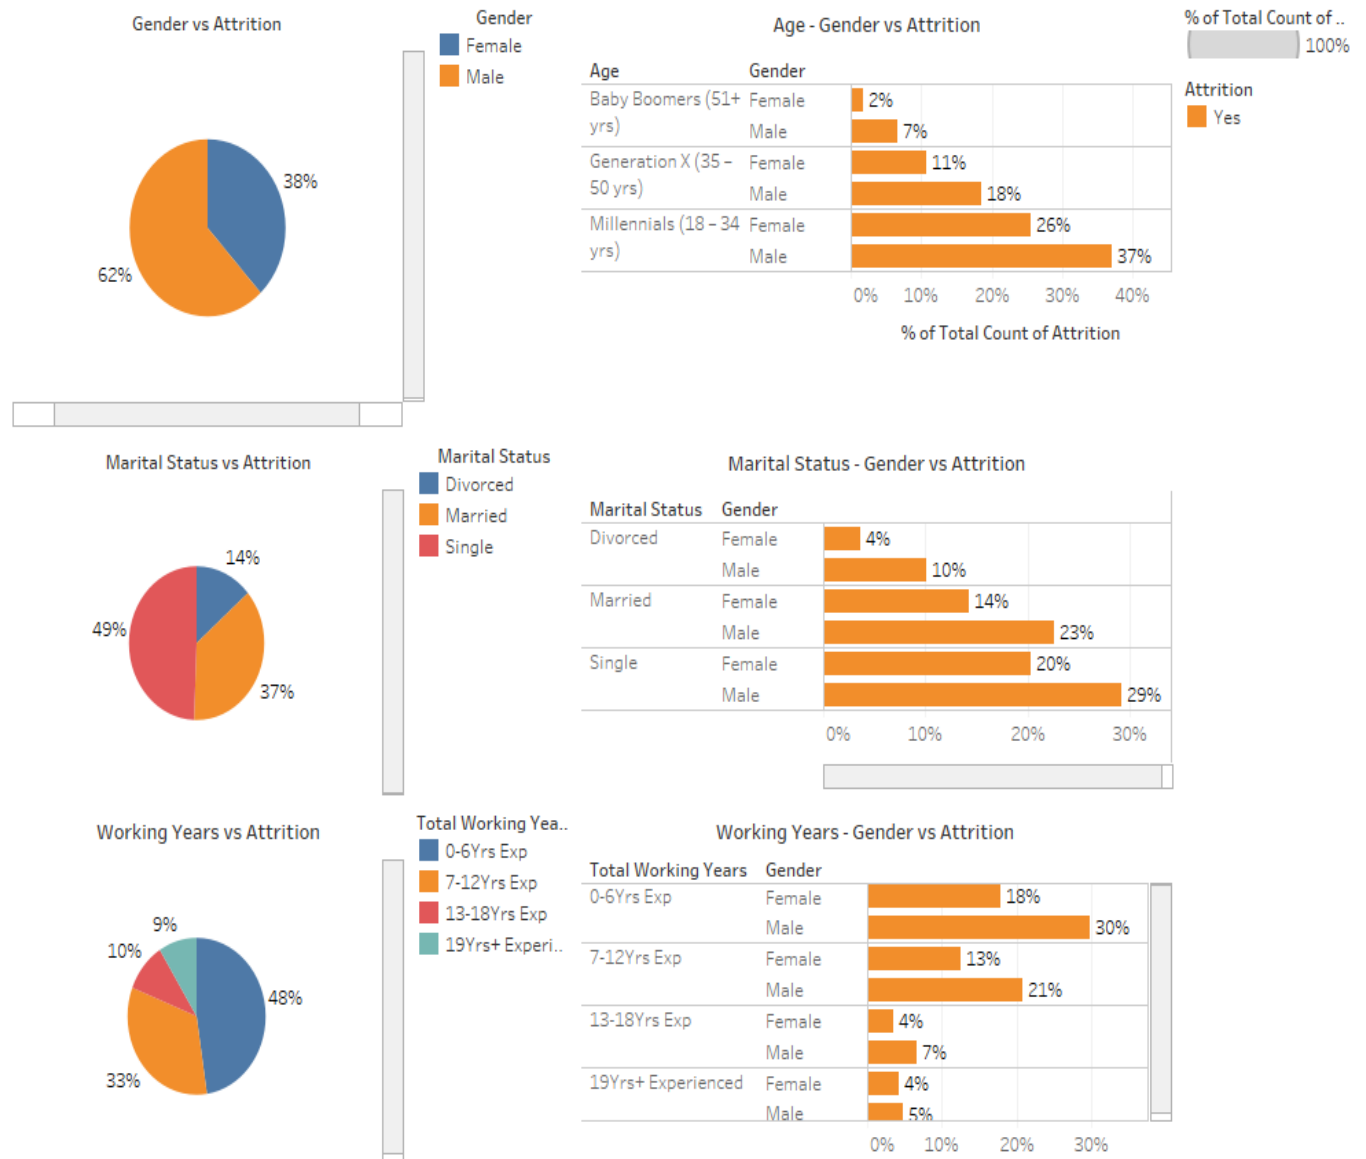
testdata



Training Data Visualization



Training Data Visualization Contd...





TECHNIQUE 1: Model Building WITHOUT any data manipulation

1. **Step 1**

First, we checked Train Data (trainData.csv) for **Blanks or '?'** (question marks) in the data. We were unable to find any.

2. **Step 2**

We then created following 4 models: [Logistic Regression](#), [Decision Tree](#), [Random Forest](#) and [Support Vector Machines](#) to predict our target Variable "Attrition" comparing with all 32 columns (variables) of the "trainData.csv".

3. **Step 3**

We then calculated the best threshold that gives the best accuracy and least total number of errors for each model. Details below:

Model#	<u>Algorithm used</u>	<u>Threshold Value</u>
MODEL 1	LOGISTIC REGRESSION (Attrition ~ .)	Any prediction value > 0.51 , "Yes", "No"
MODEL 2	DECISION TREE (Attrition ~ .) WITH INFORMATION GAIN	tuneLength = 14
MODEL 3	DECISION TREE (Attrition ~ .) WITH GINI INDEX	tuneLength = 9
MODEL 3	RANDOM FOREST (Attrition ~ .)	tuneLength = 5
MODEL 4	SVM (Attrition ~ .)	cost = 9, gamma=0.015

Confusion Matrix and accuracy of all the 4 models are shown in the next page.


4. **Step 4**

We then compared the accuracy as well as the Sensitivity and Specificity of all the 4 models to decide which model is best for predicting.

Confusion and Accuracy : Without any data manipulation

<u>LOGISTIC REGRESSION OUTPUT</u>	<u>DECISION TREE WITH GINI INDEX OUTPUT</u>
<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 360 36 Yes 12 33</p> <p>Accuracy : 0.8912 95% CI : (0.8583, 0.9186) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.0025818</p> <p>Kappa : 0.5196 McNemar's Test P-Value : 0.0009009</p> <p>Sensitivity : 0.47826 Specificity : 0.96774 Pos Pred Value : 0.73333 Neg Pred Value : 0.90909 Prevalence : 0.15646 Detection Rate : 0.07483 Detection Prevalence : 0.10204 Balanced Accuracy : 0.72300</p> <p>'Positive' Class : Yes</p>	<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 356 48 Yes 16 21</p> <p>Accuracy : 0.8549 95% CI : (0.8185, 0.8864) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.2809805</p> <p>Kappa : 0.3222 McNemar's Test P-Value : 0.0001066</p> <p>Sensitivity : 0.30435 Specificity : 0.95699 Pos Pred Value : 0.56757 Neg Pred Value : 0.88119 Prevalence : 0.15646 Detection Rate : 0.04762 Detection Prevalence : 0.08390 Balanced Accuracy : 0.63067</p> <p>'Positive' Class : Yes</p>
<u>RANDOM FOREST OUTPUT</u>	<u>SVM OUTPUT</u>
<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 359 42 Yes 13 27</p> <p>Accuracy : 0.8753 95% CI : (0.8408, 0.9046) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.0355092</p> <p>Kappa : 0.43 McNemar's Test P-Value : 0.0001597</p> <p>Sensitivity : 0.39130 Specificity : 0.96505 Pos Pred Value : 0.67500 Neg Pred Value : 0.89526 Prevalence : 0.15646 Detection Rate : 0.06122 Detection Prevalence : 0.09070 Balanced Accuracy : 0.67818</p> <p>'Positive' Class : Yes</p>	<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 360 36 Yes 12 33</p> <p>Accuracy : 0.8912 95% CI : (0.8583, 0.9186) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.0025818</p> <p>Kappa : 0.5196 McNemar's Test P-Value : 0.0009009</p> <p>Sensitivity : 0.47826 Specificity : 0.96774 Pos Pred Value : 0.73333 Neg Pred Value : 0.90909 Prevalence : 0.15646 Detection Rate : 0.07483 Detection Prevalence : 0.10204 Balanced Accuracy : 0.72300</p> <p>'Positive' Class : Yes</p>

Model Accuracy Graph : Without any data manipulation





Technique 1 : Conclusion

Top 2 Best Models as per Technique 1

Technique#	Best Model	Threshold	Accuracy	Kappa	Sensitivity	Specificity
Technique 1	Logistic Regression (var=32)	0.51	0.8912	0.5196	0.47826	0.96774
Technique 1	SVM (Var=32)	c=9,g=0.015	0.8912	0.5196	0.47826	0.96774

After comparison, Logistic Regression and SVM model has the best **accuracy of 0.8912** and **Sensitivity of 0.4783**. Hence, these two are the best models to predict data if we compare Attrition with all the 32 remaining columns and without any data manipulation.

We will use these accuracies of each model as the base to tune and create a better model.



FEATURE ENGINEERING:

Data Manipulation:

1. Removing Outliers:

What are outliers?

- In statistics, an outlier is an observation point that is distant from other observations. An outlier may be due to variability in the measurement or it may indicate experimental error; the latter are sometimes excluded from the data set.

Why outliers treatment is important?

- Outliers in data can distort predictions and affect the accuracy, if you don't detect and handle them appropriately especially in regression models. Because, it can drastically bias/change the fit estimates and predictions.

Outlier technique used:

- We used the Winsorizing technique to handle the outliers.
- Winsorizing is the transformation of statistics by limiting extreme values in the statistical data to reduce the effect of possibly spurious outliers. A typical strategy is to set all outliers to a specified percentile of the data; for example, a 90% winsorization would see all data below the 5th percentile set to the 5th percentile, and data above the 95th percentile set to the 95th percentile.



2. Scaling (Z standardization scaling)

Why Scaling?

- The purpose of scaling is to put all the results onto a common scale. This is to avoid models biasness, this helps us to ensure models are not biased based on different ranges of numeric variables, this in turn helps in increasing the accuracy of the model

Use of scale() function:

- We have used the scale() function of R to scale all numeric variables that use the Z standardization
- scale() is a generic function whose default method centers and/or scales the columns of a numeric matrix, within the range of +3 to -3.

e.g. See below the scaling for "Age" variable

Before scaling	After scaling
<pre>> summary(orgData["Age"]) Age Min. :18.00 1st Qu.:30.00 Median :36.00 Mean :36.92 3rd Qu.:43.00 Max. :60.00</pre>	<pre>> orgData["Age"] <- as.data.frame(scale(orgData["Age"])) > summary(orgData["Age"]) Age Min. :-2.0715 1st Qu.: -0.7579 Median :-0.1011 Mean : 0.0000 3rd Qu.: 0.6651 Max. : 2.5260</pre>



TECHNIQUE 2: Model Building WITH All Features & AFTER Data manipulation

1. **Step 1:**

- We again created following 5 models: Logistic Regression, Decision Tree, Random Forest and Support Vector Machines to predict our target Variable "Attrition" comparing with all 32 columns (variables) of the "trainData.csv" with outliers and scaling done

2. **Step 2:**

- We then calculated the best threshold that gives the best accuracy and least total number of errors for each model. Details below:

Model#	<u>Algorithm used</u>	<u>Threshold Value</u>
MODEL 1	LOGISTIC REGRESSION (Attrition ~ .)	Any prediction value > 0.56 , "Yes", "No"
MODEL 2	DECISION TREE (Attrition ~ .) WITH INFORMATION GAIN	tuneLength = 14
MODEL 3	DECISION TREE (Attrition ~ .) WITH GINI INDEX	tuneLength = 9
MODEL 4	RANDOM FOREST (Attrition ~ .)	tuneLength = 500
MODEL 5	SVM (Attrition ~ .)	cost = 7, gamma=0.015




Confusion and Accuracy:

All Features & AFTER Data manipulation

<u>LOGISTIC REGRESSION OUTPUT</u>	<u>DECISION TREE WITH GINI INDEX OUTPUT</u>
<p>Confusion Matrix and Statistics</p> <pre> Reference Prediction No Yes No 368 40 Yes 4 29 Accuracy : 0.9002 95% CI : (0.8684, 0.9266) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.0003614 Kappa : 0.52 McNemar's Test P-Value : 1.317e-07 Sensitivity : 0.42029 Specificity : 0.98925 Pos Pred value : 0.87879 Neg Pred value : 0.90196 Prevalence : 0.15646 Detection Rate : 0.06576 Detection Prevalence : 0.07483 Balanced Accuracy : 0.70477 'Positive' Class : Yes </pre>	<p>Confusion Matrix and Statistics</p> <pre> Reference Prediction No Yes No 356 48 Yes 16 21 Accuracy : 0.8549 95% CI : (0.8185, 0.8864) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.2809805 Kappa : 0.3222 McNemar's Test P-Value : 0.0001066 Sensitivity : 0.30435 Specificity : 0.95699 Pos Pred value : 0.56757 Neg Pred value : 0.88119 Prevalence : 0.15646 Detection Rate : 0.04762 Detection Prevalence : 0.08390 Balanced Accuracy : 0.63067 'Positive' Class : Yes </pre>
<u>RANDOM FOREST OUTPUT</u>	<u>SVM OUTPUT</u>
<p>Confusion Matrix and Statistics</p> <pre> Reference Prediction No Yes No 370 55 Yes 2 14 Accuracy : 0.8707 95% CI : (0.8358, 0.9006) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.06318 Kappa : 0.2874 McNemar's Test P-Value : 5.675e-12 Sensitivity : 0.20290 Specificity : 0.99462 Pos Pred value : 0.87500 Neg Pred value : 0.87059 Prevalence : 0.15646 Detection Rate : 0.03175 Detection Prevalence : 0.03628 Balanced Accuracy : 0.59876 'Positive' Class : Yes </pre>	<p>Confusion Matrix and Statistics</p> <pre> Reference Prediction No Yes No 361 38 Yes 11 31 Accuracy : 0.8889 95% CI : (0.8558, 0.9167) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.0039921 Kappa : 0.4993 McNemar's Test P-Value : 0.0002038 Sensitivity : 0.44928 Specificity : 0.97043 Pos Pred value : 0.73810 Neg Pred value : 0.90476 Prevalence : 0.15646 Detection Rate : 0.07029 Detection Prevalence : 0.09524 Balanced Accuracy : 0.70985 'Positive' Class : Yes </pre>

Model Accuracy Graph: All Features & AFTER Data manipulation





Technique 2 : Conclusion

Top 2 Best Models as per Technique 2:

Technique#	Best Model	Threshold	Accuracy	Kappa	Sensitivity	Specificity
Technique 2	Logistic Regression (var=32)	0.56	0.9002	0.5200	0.4203	0.98925
Technique 2	SVM (var=32)	c=7, g=0.015	0.8889	0.4993	0.449275	0.97043

After, implementing Technique 2, now when we compare all the models we can conclude that the Logistic Regression model is giving the best **accuracy of 0.9002** among all.

If we compare this Logistic model with the previous one (Technique 1) then we see that the accuracy has improved marginally from **0.8912** to **0.9002** (increase of **0.009**) but there is a slight decrease of **0.0580** (approx. 5.8%) in Sensitivity (from **0.47826** to **0.42029**)


TECHNIQUE 3:

Model Building AFTER Data manipulation including only Important features


Important feature selection

- We used the VarImp() function in R for feature selection for Decision Tree, Random Forest and SVM.
- For Logistic we kept on iterating the model and eliminating the variables with high p-values until we had only features which had p-values greater than 0.05 or up to 0.1 if it affects the overall accuracy.

16 Important features as per Decision Tree model with Information Gain:




19 Important features as per Decision Tree model with Gini Index:




Important features as per Random forest model (acc vs Gini):

We used only variables with value greater than 0 for MeanDecreaseAccuracy and Variables with value greater than 9 for MeanDecreaseGini for our models.



Important features as per SVM model:

We used only variables with value greater than 0.01 for our model)





MODEL CREATION USING IMPORTANT VARIABLES

Below are the Models created selecting only important features:


Model#	<u>Algorithm used</u>	<u>Variable Count</u>	<u>Threshold Value</u>
MODEL 1	LOGISTIC REGRESSION	18	Any prediction value > 0.53 , "Yes", "No"
MODEL 2	LOGISTIC REGRESSION - chisq	20	Any prediction value > 0.48 , "Yes", "No"
MODEL 3	DECISION TREE WITH INFORMATION GAIN	16	tuneLength = 12
MODEL 4	DECISION TREE WITH GINI INDEX	19	tuneLength = 9
MODEL 5	RANDOM FOREST (Gini > 9)	14	ntree = 55
MODEL 6	SVM	21	cost = 20, gamma=0.005

Confusion and Accuracy: AFTER Data manipulation including only Important features

<u>LOGISTIC REGRESSION WITH ANOVA OUTPUT</u>	<u>LOGISTIC REGRESSION WITH CHISO OUTPUT</u>
<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 365 49 Yes 7 20</p> <p>Accuracy : 0.873 95% CI : (0.8383, 0.9026) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.04779</p> <p>Kappa : 0.3604 McNemar's Test P-Value : 4.281e-08</p> <p>Sensitivity : 0.28986 Specificity : 0.98118 Pos Pred Value : 0.74074 Neg Pred Value : 0.88164 Prevalence : 0.15646 Detection Rate : 0.04535 Detection Prevalence : 0.06122 Balanced Accuracy : 0.63552</p> <p>'Positive' Class : Yes</p>	<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 363 42 Yes 9 27</p> <p>Accuracy : 0.8844 95% CI : (0.8508, 0.9127) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.00896</p> <p>Kappa : 0.4559 McNemar's Test P-Value : 7.433e-06</p> <p>Sensitivity : 0.39130 Specificity : 0.97581 Pos Pred Value : 0.75000 Neg Pred Value : 0.89630 Prevalence : 0.15646 Detection Rate : 0.06122 Detection Prevalence : 0.08163 Balanced Accuracy : 0.68356</p> <p>'Positive' Class : Yes</p>
<u>DECISION TREE WITH INFO GAIN OUTPUT</u>	<u>DECISION TREE WITH GINI INDEX OUTPUT</u>
<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 357 52 Yes 15 17</p> <p>Accuracy : 0.8481 95% CI : (0.8111, 0.8803) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.4278</p> <p>Kappa : 0.2636 McNemar's Test P-Value : 1.092e-05</p> <p>Sensitivity : 0.24638 Specificity : 0.95968 Pos Pred Value : 0.53125 Neg Pred Value : 0.87286 Prevalence : 0.15646 Detection Rate : 0.03855 Detection Prevalence : 0.07256 Balanced Accuracy : 0.60303</p> <p>'Positive' Class : Yes</p>	<p>Confusion Matrix and Statistics</p> <p>Reference Prediction No Yes No 356 48 Yes 16 21</p> <p>Accuracy : 0.8549 95% CI : (0.8185, 0.8864) No Information Rate : 0.8435 P-Value [Acc > NIR] : 0.2809805</p> <p>Kappa : 0.3222 McNemar's Test P-Value : 0.0001066</p> <p>Sensitivity : 0.30435 Specificity : 0.95699 Pos Pred Value : 0.56757 Neg Pred Value : 0.88119 Prevalence : 0.15646 Detection Rate : 0.04762 Detection Prevalence : 0.08390 Balanced Accuracy : 0.63067</p> <p>'Positive' Class : Yes</p>

RANDOM FOREST Gini > 9	SVM
<p>Confusion Matrix and Statistics</p> <pre> Reference Prediction No Yes No 365 45 Yes 7 24 </pre> <p> Accuracy : 0.8821 95% CI : (0.8483, 0.9107) No Information Rate : 0.8435 P-value [Acc > NIR] : 0.01302 Kappa : 0.4241 McNemar's Test P-value : 2.882e-07 Sensitivity : 0.34783 Specificity : 0.98118 Pos Pred Value : 0.77419 Neg Pred Value : 0.89024 Prevalence : 0.15646 Detection Rate : 0.05442 Detection Prevalence : 0.07029 Balanced Accuracy : 0.66450 'Positive' Class : Yes </p>	<p>Confusion Matrix and Statistics</p> <pre> Reference Prediction No Yes No 367 47 Yes 5 22 </pre> <p> Accuracy : 0.8821 95% CI : (0.8483, 0.9107) No Information Rate : 0.8435 P-value [Acc > NIR] : 0.01302 Kappa : 0.4061 McNemar's Test P-value : 1.303e-08 Sensitivity : 0.31884 Specificity : 0.98656 Pos Pred Value : 0.81481 Neg Pred Value : 0.88647 Prevalence : 0.15646 Detection Rate : 0.04989 Detection Prevalence : 0.06122 Balanced Accuracy : 0.65270 'Positive' Class : Yes </p>

Model Accuracy Graph: **AFTER Data manipulation including only important features**



**Conclusion:**

Top 2 Best Models as per Technique 3

Technique#	Best Model	Threshold	Accuracy	Kappa	Sensitivity	Specificity
Technique 3	Logistic Regression (var=20)	0.48	0.8844	0.45591	0.39130	0.97581
Technique 3	Random Forest (gini>9) (var=14)	ntree=55	0.8821	0.4241	0.34783	0.98118

After, implementing Technique 3, now when we compare all the models we can conclude that the Logistic Regression model is giving the best **accuracy of 0.8844** among all.


Final Conclusion:

Objective 1: Top 10 important features for attrition:

Below are the top 10 important features from the output using the VarImp function on the various models:


1. Overtime
2. Number of Companies worked
3. Environmental Satisfaction
4. Distance from Home
5. Job Satisfaction
6. Job Involvement
7. Total Working Years
8. Work Life Balance
9. Business Travel
10. Age

Find below the charts showing the relationship between the Top 10 features and Attrition:






Working Years vs Attrition



Total Working Yea..



Work life Balance vs Attrition



Work Life Balance




Count of Attrition

168

% of Total Count of Attriti..

100%


Travel vs Attrition




Business Travel



Age vs Attrition



Age





OBJECTIVE 2:

Creating the best predictive model:

Below is the list of all the Best models based on Techniques 1 to 3

Technique#	Best Model	Threshold	Accuracy	Kappa	Sensitivity	Specificity
Technique 1	Logistic Regression (var=32)	0.51	0.8912	0.5196	0.47826	0.96774
Technique 1	SVM (var=32)	c=9, g=0.015	0.8912	0.5196	0.47826	0.96774
Technique 2	Logistic Regression (var=32)	0.56	0.9002	0.5200	0.42029	0.98924
Technique 3	Logistic Regression (var=20)	0.48	0.8844	0.45591	0.39130	0.97581

If you have a look at the above table then you can see that we were able to improve the accuracy as well as the Kappa Value when we created Logistic models comparing all features and after imputing the outliers as well as scaling (Technique 2) when compared to Technique 1

Best Model after all 3 Techniques is:


Technique#	Best Model	Threshold	Accuracy	Kappa	Sensitivity	Specificity
Technique 2	Logistic Regression (var=32)	0.56	0.9002	0.5200	0.42029	0.98924

Significant findings from Best Logistic Regression Model:

Based on the Logistic Model Predicting we see that we were able to make 29 correct predictions which is 42% when compared to expected prediction. Summary of the findings from predictions is mentioned below:

1. Possibility of Attrition is high in **Males** than Females.
2. **Sales Department** has the highest Attrition rate.
3. Possibility of Attrition is high in **One-Three Year** old employees.
4. **Sad** employees (who rated 1 or 2 for all the three satisfaction parameters: Job Satisfaction, Environment Satisfaction and Relationship Satisfaction) are more likely to leave the organization.

Significance Attrition findings from Logistic Predictive Model





**A BIG THANK YOU TO ALL
THE TEACHERS
AND SUPPORT STAFF
OF IMARTICUS LEARNING**

